

Lecture Notes 3

Partitioning (Sharding)

- Split data over multiple nodes
- Partitioning improves performance

Replication

Consistency

- All nodes see same data at same time
- Consistency - there are multiple types
 - Strong - offers up to date data but at cost of high latency
 - Eventual - offers low latency but may reply to read requests with stale data since all nodes of the database may not have the updated data

Caching

- Data is written once and read many times - many systems have data with read-write ratio > 1
- Cache Consistency Problems
 - Write through, Write back
- Cache replacement policies

Caching in Distributed Systems

CAP Theorem

- Consistency, Availability, and Partition Tolerance
- Consistency - all nodes see same data at same time
- Availability - node failures do not prevent survivors from continuing to operate
- Partition Tolerance - the system continues to operate despite message loss due to network failures

- CP Systems - Incorporates network partitions into their failure model and Strong Consistency, ie Twitter Manhattan Storage System
- AP Systems - Eventual consistency and highly available, ie most NoSQL databases, HBase

NoSQL

- NoSQL: Not Only SQL
 - HBase, Cassandra, MongoDB, CouchDB
 - Schema not predefined
 - For Horizontal Scaling
- Modern applications create massive volumes of semi-structured data
- Scale out architectures to reach the demand
- Traditional RDBMS were not designed to cope with the scale that faces modern applications

RDBMS vs NoSQL

- BASE, not ACID
 - RDBMS (ACID) - Atomicity consistency Isolation Durability
 - NoSQL (BASE) - Basically Available Soft state Eventual consistency
- NoSQL: Schema-free, easy replication support, simple API, and High Availability (eventual consistent)
- What is missing in NoSQL?
 - No joins support in the Database - Applications need to take care of joins
 - No complex transactions as in SQL
 - No constraints support (values could be null or empty)
- NoSQL is Highly Available (eventual consistent) and Scalable

Zookeeper

- Coordination in large scale distributed systems
 - locking service (paxos)
 - dynamic configuration management
 - group membership, etc.
- Zookeeper is already used in HBase, HDFS
- Design goals - High availability, fault tolerance, performance
- Runs on a collection of machines and is designed to be highly available, so applications can depend on it
- Zookeeper can help you avoid introducing single points of failure into your system, so you can build a

reliable application

Interesting Distributed System Design problems

- Tiny URL
- Google search
- Facebook Newsfeed
- Dropbox
- Twitter

Quiz on Friday

- There is an exam (1:10 to 7:45)
- 1 sentence answer, true false, 1 design (Interesting distributed system design questions)